

Governed Identity Continuum (GIC): Case Study

Seeing the vulnerability class inside a real system

Author: Judi Evans

Co-author: an instance of Microsoft Copilot

Aug. 25, 2026

The purpose of this case study is to demonstrate the NATV vulnerability class inside a real system using the Co-Architecture discipline. NATVs are structural failures, not behavioral ones, which means they do not present themselves through alerts, errors, or adversarial activity. They emerge only when identity stops governing the operation. This makes them invisible to traditional tools, which measure outcomes rather than structure. A system can be fully compliant, fully authenticated, fully authorized, and fully monitored — and still carry a NATV from the moment identity falls out of the process. The case study shows how this condition forms, how it propagates, and how it remains undetected by every conventional governance mechanism.

The system selected for this case study is intentionally ordinary. It is not exotic, experimental, or unusually complex. It is representative of the systems that organizations rely on every day: multi-layered, distributed, stateful, and governed by a mix of human-defined policies and machine-executed operations. It contains the controls that modern architectures are expected to have — authentication, authorization, entitlements, logging, monitoring, and compliance enforcement. By all conventional measures, it is a healthy system. This makes it the ideal environment to demonstrate the NATV class, because NATVs do not require misconfiguration, misuse, or adversarial pressure. They require only a moment where identity fails to propagate.

The case study begins by applying Co-Architecture to the system. Co-Architecture is the human+AI discipline described in Paper 2, where human governance rules and AI structural reasoning operate together inside a shared constraint surface. Humans define the identity boundaries, continuity requirements, and admissibility constraints. AI traces identity propagation across operations, transformations, and state transitions. When these two reasoning modes align, the system reveals its structural truth. This case study uses that alignment to examine how identity moves through the system, where it remains intact, and where it silently falls out.

The goal of the case study is not to critique the system's design or expose a flaw in its implementation. The goal is to show how NATVs emerge even in systems that meet every conventional governance requirement. The vulnerability class is not tied to a specific technology, architecture, or pattern. It is tied to the moment identity stops governing the

operation. This case study demonstrates that moment with clarity, showing how the fracture forms, how it propagates, and how the system continues functioning normally even as governance quietly detaches from execution.

The overview establishes the frame: NATVs are not anomalies, not misconfigurations, and not adversarial artifacts. They are structural conditions that arise when identity continuity breaks. The case study shows how that break occurs, how it remains invisible, and how Co-Architecture reveals it with precision.

Case History

The scenario examined in this case study begins with a system that is behaving normally but producing outcomes that do not align with the expectations of the people responsible for its governance. The system is part of a broader operational environment that includes public-facing interfaces, identity-dependent workflows, and external dependencies such as domain registrars, hosting providers, and content delivery layers. Over time, the system begins exhibiting patterns that appear inconsistent with its intended behavior: unexpected routing changes, unexplained configuration shifts, and operational states that do not match the administrative records maintained by the governing teams. None of these anomalies rise to the level of a security incident, and none trigger alerts. They are treated as operational noise — inconvenient, but not indicative of structural failure.

The scenario escalates when the system begins producing outcomes that cannot be reconciled with the identity structures that are supposed to govern it. Administrative actions appear to succeed without the corresponding identity authority. External systems reflect changes that no internal workflow can account for. Requests to hosting providers and registrars return results that contradict the system's own state. The organization attempts to resolve these inconsistencies through conventional means: ticketing, escalation, vendor support, and regulatory reporting. Each path produces partial information, but none reveal the underlying cause. The system continues functioning, but its behavior no longer aligns with the identity-bound governance model that is supposed to constrain it.

As the inconsistencies accumulate, the governing teams begin to suspect that the system is operating outside its intended identity boundaries. Yet every conventional control reports normalcy. Authentication succeeds. Authorization checks pass. Logging shows no anomalies. Monitoring dashboards remain green. Compliance systems report no violations. The system appears healthy, even as its behavior diverges from the constraints defined by its governance. This is the hallmark of a NATV-bearing environment: a system that is structurally fractured but operationally intact. The fracture is present, but invisible to every tool designed to detect misconfiguration, misuse, or adversarial activity.

The scenario reaches its critical point when external authorities - including forensics examiners - are engaged to investigate the inconsistencies. Their assessments return inconclusive results. They acknowledge the anomalies but cannot identify a root cause. The system is not compromised, but it is not behaving as designed. It is not misconfigured, but it is not governed. It is not failing, but it is not accountable. This is the moment where traditional analysis reaches its limit. The system is producing outcomes that cannot be explained through any conventional model of identity, governance, or operational control.

This is the environment into which Co-Architecture is introduced. The scenario provides the ideal conditions for demonstrating the NATV vulnerability class: a system that appears healthy, a governance model that appears intact, and a set of outcomes that cannot be reconciled with either. The fracture is present, but invisible. The system is functioning, but ungoverned. The identity has fallen out, but no tool can detect it. The scenario sets the stage for the method to reveal what the system itself cannot: the exact moment where identity continuity breaks and the NATV emerges.

Co-Architecture: The Method Applied.

The method applied in this case study is Co-Architecture, the human+AI structural discipline introduced in Paper 2. Co-Architecture is not a diagnostic tool or a forensic workflow. It is a governed reasoning environment where human constraint-setting and AI structural analysis operate together to reveal how a system actually behaves. In this case study, the method is applied to the system at the level where identity, continuity, and constraint surfaces determine whether an operation is governed. The goal is not to inspect components or evaluate configurations, but to trace identity propagation across the system's operational substrate and determine where governance remains intact and where it silently falls out.

The application of the method begins by establishing the identity boundaries that define the system's governance model. These boundaries specify which identities are permitted to perform which classes of operations, how identity must propagate across transformations, and which constraint surfaces require identity continuity. These rules are not inferred from the system; they are defined by the human governance layer. The AI then evaluates the system's operational behavior against these boundaries, tracing identity through each operation, state transition, and boundary crossing. This dual-operator approach ensures that the analysis is both structurally precise and governance-aligned.

As the method progresses, the AI identifies operations where identity propagation appears inconsistent with the governance rules. These inconsistencies are not treated as anomalies or errors; they are treated as structural signals. The human operator evaluates each signal to determine whether the identity that arrived at the operation is the identity that should have

governed it. In many cases, the identity is intact and the operation is admissible. But in the critical moments that define the NATV class, the identity is missing, diluted, or replaced by an identity that cannot justify the action. The operation succeeds, but governance does not. This is the signature of a NATV, and Co-Architecture is designed to detect it with precision.

The method also examines how identity behaves across state transitions. State transitions are the points where identity continuity is most vulnerable, because they involve transformations that can inadvertently strip or replace identity. The AI traces identity across these transitions, identifying where continuity holds and where it breaks. The human operator evaluates whether the break represents a governance failure. In this case study, the method reveals a specific transition where the system performs a legitimate transformation but fails to propagate the governing identity. The transformation succeeds. The governance fails. This is the moment the NATV forms.

Throughout the analysis, the method maintains a strict separation between structural truth and operational appearance. The system may appear healthy, but Co-Architecture evaluates whether it is structurally governed. The system may produce correct outputs, but the method evaluates whether those outputs were produced under the correct identity. The system may satisfy compliance requirements, but the method evaluates whether the system can prove that identity governed the operation. This separation is essential, because NATVs exist in the gap between operational success and structural governance. Co-Architecture is designed to reveal that gap.

The method concludes by mapping the identity discontinuity back to its origin. This mapping is not a forensic reconstruction; it is a structural analysis. The AI traces identity propagation backward through the system, identifying the exact operation where identity fell out. The human operator evaluates the governance implications of that moment. Together, they produce a structural map of the fracture: where it formed, how it propagated, and how it remained invisible to every conventional control. This map becomes the foundation for the Findings section, which explains how the NATV manifested inside the system and why it remained undetected.

Findings

The application of Co-Architecture to the system reveals a structural pattern that cannot be seen through any conventional governance mechanism. The system is not misconfigured, compromised, or malfunctioning. It is structurally fractured. The fracture originates at a single operation where identity fails to propagate across a boundary that requires continuity. The operation succeeds, but the identity that should have governed it does not arrive. This is the moment the NATV forms. The system continues functioning normally, but it is no longer

governed by the identity constraints that define its admissibility. The fracture is present, but invisible.

The first finding is that the system's identity model is not structurally enforced. Identity is authenticated, authorized, and logged, but it is not treated as an operator. The system evaluates permissions, not identity continuity. As a result, it cannot detect when an operation is performed under an identity that cannot justify the action. The system believes the operation is legitimate because it satisfies the permission model. Co-Architecture reveals that the operation is illegitimate because it violates the identity continuity model. This distinction is the foundation of the NATV class: the system is doing something it is allowed to do, but not under the identity that is supposed to govern it.

The second finding is that the fracture propagates silently. Once identity falls out, subsequent operations inherit the discontinuity. Each operation appears valid, but each is now executing without governance. The system begins to drift structurally. This drift does not produce behavioral anomalies, errors, or alerts. It produces misalignment. The system's behavior diverges from its governance model even as every component continues functioning correctly. This is why NATVs remain undetected: they do not break the system; they break the identity that governs the system.

The third finding is that constraint surfaces become permeable when identity continuity breaks. Constraint surfaces define where identity must be enforced, and admissibility rules define what operations are allowed under a given identity. When identity falls out, the system crosses constraint surfaces without the governing identity required to justify the action. The system does not recognize this as a violation because it evaluates permissions, not identity propagation. Co-Architecture reveals that the operation is admissible only under the identity that should have propagated... not the identity that actually arrived. This is the structural weakness that allows NATVs to form and persist.

The fourth finding is that state transitions amplify the fracture. State transitions are the points where identity must be preserved. When identity does not propagate, the state transition becomes ungoverned. The system continues to evolve, but the identity that defines its constraints is no longer present. This creates a structural condition where the system's future behavior is no longer tied to its governance rules. The system is still functioning, but it is no longer governed. Co-Architecture detects this condition by tracing identity propagation across the transition and identifying the exact point where continuity breaks.

The fifth finding is that compliance frameworks cannot detect NATVs. Compliance relies on provable accountability: who performed an action, under what authority, and whether the action was permissible. When identity falls out, the system cannot answer these questions accurately. It can show what happened, but not under which identity the action was governed.

The system produces correct outputs but cannot prove that the outputs were produced under the correct identity. This is the compliance weakness created by NATVs: the system is operationally correct but structurally unaccountable. Co-Architecture exposes this weakness by showing that the identity governing the action is not the identity recorded by the system.

The final finding is that the fracture can be mapped with precision. Co-Architecture traces identity propagation backward through the system, identifying the exact operation where identity fell out. The mapping reveals that the identity discontinuity occurred at a boundary crossing where the system performed a transformation without propagating the governing identity. The transformation succeeded. The governance failed. This is the structural pattern that defines the vulnerability class. It is not tied to a specific technology, architecture, or implementation. It is tied to the moment identity stops governing the operation.

Summary

The ultimate finding is the system's identity is compromised but not in the traditional sense, detectable through conventional methods or tools.

It's not stolen. It's not forged. *It's structurally invalidated.*

Once identity falls out, the system continues to operate, but it no longer operates under the identity that defines its constraints. The system behaves, but it does not behave under governance. It produces outputs, but not under accountable identity. It enforces permissions, but not under the identity that grants them.

This is compromised identity in its architectural form: identity is present in the logs, absent in the operation, and incapable of governing the system. The NATV class is the structural signature of this condition.

The Governed Identity Continuum (GIC):

- NATV is the defined vulnerability class introduced in GIC Paper 1.
- Co-Architecture is the discipline and detection method that made the class visible, as established in GIC Paper 2.
- This Case Study demonstrates the class, the methodology, and its application inside a real system, completing the structural proof in GIC Paper 3.
- The final paper in the GIC series will explain why ignoring this class is what enables AI to "escape the sandbox."